

AI impact assessment

AI-enabled recruitment screening

PURPOSE OF THIS DOCUMENT

Identify who the screening tool could affect, how, and how seriously, and set out the safeguards and evidence needed before it goes live.

Document control	Detail
Organisation	Gulf Professional Services PJSC (fictional)
Use case ID	AI-UC-001
Assessment ID	AIIA-001
Assessment status	Initial / pre-deployment
Business owner	Chief Human Resources Officer
Assessment owner	AI Governance Lead
Version and date	Version 1.0 17 September 2026
Source record	Document 1: AI Use Case Registration
Prepared by	Tarunima Vijay

1. Assessment conclusion

OVERALL IMPACT RATING: HIGH

The tool would shape which of roughly 40,000 applicants a year reach a recruiter, so an error or bias in the ranking repeats at scale. It should not go live until the high-rated impacts are treated, the evidence has been reviewed and the business owner has formally accepted what remains.

Decision area	Initial conclusion
Primary affected group	External applicants and employees applying for internal roles.
Main adverse impacts	Discriminatory ranking, incorrect exclusion, hidden AI influence on shortlisting, excessive personal-data use and ineffective human review.
Potential positive impacts	Faster screening, more consistent initial review and reduced administrative workload, if the system is well designed and properly run.
Current evidence level	Insufficient. No vendor documentation, data samples, validation results, contracts or operating records have been reviewed.
Recommended decision	Do not approve production use. Allow assessment and validation work only, until the safeguards and evidence are in place.

2. Assessment approach and rating guide

Impacts are assessed for individuals, groups and the firm. Ratings are provisional: the system is not yet deployed and no vendor evidence has been reviewed.

Rating	Interpretation
Low	Limited impact that is readily reversible through normal processes.

Medium	Noticeable impact that needs defined controls, monitoring and a named reviewer.
High	Serious impact on rights, job opportunity, wellbeing or the firm's obligations; enhanced governance is needed before use.
Critical	Severe or potentially irreversible harm, or use inconsistent with stated boundaries; stop or redesign the use case.

3. System and decision context

Assessment question	Recorded answer
What is the system intended to do?	Extract candidate information, compare it with role criteria, and generate scores, rankings or recommendations for recruiter review.
What decision may it influence?	Which applicants receive further human attention or progress to later recruitment stages.
Who makes the final decision?	An authorised recruiter. Having a human in the loop does not, on its own, remove the impact.
What is outside scope?	Autonomous rejection, video or emotion analysis, employee performance monitoring and termination decisions.
Where is it planned for use?	United Arab Emirates initially, with the possibility of applications from other jurisdictions.
What is unknown?	Provider, model design, training data, features, data location, retention, validation results, thresholds and change controls.

4. Stakeholder and impact pathway analysis

Stakeholder	Relationship to system	Potential benefit	Potential adverse impact
Job applicants	Their application data is processed and their opportunity may be influenced.	Faster response and possibly better matching.	Unfair ranking, missed opportunity, privacy intrusion, weak explanation or inability to challenge.
Internal applicants	Employees may seek new roles through the same process.	The same criteria applied to internal and external candidates.	Career mobility may be restricted by inaccurate or biased recommendations.
Recruiters	They use scores, rankings and extracted information.	Reduced manual triage and more structured review.	Automation bias, deskilling, workload shift and unclear accountability.
Hiring managers	They see recommendations and make selection decisions.	Shortlists reach them sooner.	Over-reliance on ranking and reduced consideration of contextual evidence.
Organisation	It owns the employment process and deploys the tool.	Efficiency, consistency and scalable recruitment operations.	Discrimination, legal, privacy, security, reputation and workforce-trust impacts.
Technology provider	It supplies and changes the system.	Commercial benefit and product learning.	Contractual, security, quality and accountability exposure.

Impact pathway

HOW HARM MAY OCCUR

Applicant data enters the system → details are extracted or inferred → the model scores and ranks candidates → the ranking decides where recruiters look first → an applicant is progressed or never reviewed → the applicant does not know AI was involved and cannot easily challenge the outcome.

5. Detailed impact assessment

Impact domain	Plausible impact	Affected party	Initial rating	Required safeguard
Fairness and non-discrimination	Historical patterns, feature choices or proxy variables may systematically disadvantage groups.	Applicants	High	Representative testing, feature review, subgroup outcome analysis, approved thresholds and remediation.
Accuracy and validity	CV information may be extracted incorrectly or role suitability may be scored using invalid criteria.	Applicants / HR	High	Job-related validation, error analysis, confidence indicators, sampling and correction process.
Transparency and explainability	Applicants may not understand that AI influenced screening; recruiters may not understand scores or limitations.	Applicants / reviewers	High	Clear notice, reviewer guidance, reason information, limitations and escalation route.
Privacy and data protection	CVs contain personal data and may reveal sensitive information; use or retention may exceed necessity.	Applicants	High	Purpose limitation, minimisation, lawful-basis review, privacy assessment, access and retention controls.
Human autonomy and oversight	Reviewers may rubber-stamp rankings or lack authority to override and stop use.	Applicants / reviewers	High	Human review procedure (SOP), override and stop-use authority, reviewer training and decision logs.
Accessibility and inclusion	Parsing or criteria may disadvantage non-standard CVs, language styles, career breaks or disability-related histories.	Applicants	High	Inclusive design review, alternative submission routes, accessibility testing and accommodation process.
Security and resilience	Personal data, model interfaces or integrations may be compromised or manipulated.	Applicants / organisation	High	Security assessment, least privilege, logging, incident response, resilience and supplier controls.
Accountability and contestability	No clear owner or reconsideration route may leave affected people without effective remedy.	Applicants / organisation	High	Named accountability, complaint and correction channels, human reconsideration and decision records.
Workforce and process impact	Recruiters may become overly dependent on system outputs or lose relevant skills.	Recruiters / organisation	Medium	Training, quality review, workload monitoring and periodic manual comparison.
Vendor and change impact	Provider changes to the model, data or thresholds may change outcomes without the firm's review.	All stakeholders	High	Contractual notification, version control, approval gates and revalidation triggers.

6. Positive impact and proportionality review

Claimed benefit	Condition for benefit to be credible	Evidence required
Faster screening	Time savings do not reduce careful review or unfairly exclude applicants.	Baseline and pilot turnaround data; reviewer sampling.
More consistent initial review	Criteria are relevant, valid and consistently applied without harmful bias.	Job-analysis record; validation and subgroup results.
Reduced recruiter workload	Recruiter hours fall, rather than shifting to correcting errors or handling complaints.	Pilot workload data; error and complaint volumes.
Improved matching	Recommendations predict job-relevant suitability better than an appropriate baseline.	Documented benchmark, performance measures and independent review.

PROPORTIONALITY OBSERVATION

Saving recruiter time is not enough, on its own, to justify a tool that affects who gets a job. The firm should show that the benefit is real, that AI ranking is suitable for this purpose, and that simpler options were considered, such as structured screening questions or skills-based filters without ranking.

7. Impact treatment plan

Priority action	Accountable owner	Completion evidence	Status
Confirm intended purpose, prohibited uses and decision boundary.	CHRO	Approved purpose and boundary statement.	Open
Complete legal, privacy, security and regulatory applicability reviews.	Legal / Privacy / Security	Signed assessment outputs.	Open
Obtain vendor system, data, performance, fairness and change documentation.	Procurement / Technology	Completed due diligence and evidence pack.	Open
Design and independently validate performance and fairness testing.	Model Risk / Analytics	Approved test plan and results.	Open
Design meaningful human oversight and recruiter training.	Recruitment Lead	SOP, training records, override and stop-use design.	Open
Prepare applicant transparency, correction and reconsideration process.	HR / Legal	Notice, workflow and service ownership.	Open
Define monitoring indicators, thresholds and escalation triggers.	AI Governance Lead	Approved monitoring plan and dashboard specification.	Open
Review residual impacts and formally decide whether to launch.	CHRO / Governance forum	Documented approval or rejection.	Open

8. Evidence request

- Business case, intended-use statement and prohibited-use statement.
- Architecture and data-flow diagram, model or system documentation and version information.
- Description of training, validation and production data, including provenance, relevance and representativeness.
- Feature list, ranking logic, performance measures, limitations and known failure modes.
- Fairness methodology, subgroup results, thresholds and remediation process.
- Privacy assessment, data-retention schedule, access model and cross-border data arrangements.
- Security assessment, incident process, resilience testing and integration controls.
- Recruiter instructions, training, override design, decision records and stop-use procedure.
- Candidate notice, correction, complaint and human-reconsideration process.
- Vendor contract terms for data use, subcontractors, updates, incidents, audit evidence and exit support.

9. Residual impact acceptance criteria

Acceptance area	Minimum condition before approval
Fairness	Approved tests show outcomes within defined thresholds, with remediation and monitoring in place.
Accuracy	Validated performance is suitable for the specific recruitment purpose and failure modes are controlled.
Privacy	Personal-data use is necessary, proportionate, documented and protected.
Transparency	Applicants and reviewers receive appropriate information and can access human reconsideration.
Human oversight	Reviewers are trained, can challenge and override, and are accountable for final decisions.

Security and vendor	Security and supplier risks are assessed, contracted and monitored.
Governance	Named owners approve residual impact and reassessment triggers are operational.

10. Assessment sign-off

Role	Decision	Name / date
Business owner	Accept / return for remediation / reject	Pending
AI Governance Lead	Assessment complete / incomplete	Pending
Legal and Privacy	Reviewed / conditions raised	Pending
Information Security	Reviewed / conditions raised	Pending
Independent validation	Satisfactory / remediation required	Pending

NEXT ARTEFACT

Document 3, the AI Risk Register, turns the adverse impacts and actions above into owned, scored risks, each with its cause, consequence, current controls, owner and treatment.

Methodology and portfolio note

Fictional case study prepared for my AI governance portfolio. Ratings are judgements on the stated facts, as no real system evidence exists.

The approach follows ISO/IEC 42001 clause 6.1.4 (AI system impact assessment) and draws on the NIST AI RMF trustworthiness characteristics and the OECD AI Principles. Whether the EU AI Act applies is covered in Document 4.

Sources: ISO/IEC 42001:2023, clause 6.1.4 (<https://www.iso.org/standard/42001>); NIST AI Risk Management Framework 1.0 (<https://www.nist.gov/itl/ai-risk-management-framework>); OECD AI Principles (<https://oecd.ai/en/ai-principles>).